

PAPRICA - PATHway PRediction by phylogenetic plAcement (formerly Genome Finder)

Please contact bowmanjs@ldeo.columbia.edu with any questions.

Quick Start

Download, install, and test the installation of the dependencies for PAPRICA_LIGHT as noted under the requirements section. You will need to add the pplacer, mothur, and pathway-tools directories to your PATH and re-source the appropriate startup file (e.g. .bash_profile, .profile, .bashrc). All of these are mainstream software with pre-compiled executables for OSX and Linux (ppplacer and mothur) or installers (pathway-tools). Pathway-tools requires that you request a license before downloading (free for academic use), they are fairly quick to provide this.

From your home directory (or wherever you like to install things) download the scripts using git with:

```
git clone https://github.com/bowmanjeffs/genome_finder.git
```

This will generate a directory called genome_finder. OPEN EACH of the Python scripts and look for the “user settable variables” near the top. Modify them according to your directory structure. The PAPRICA database name should remain the same, but you will need to point to the right location. If you’ve downloaded everything to your home directory, and installed Pathway-Tools to its default location, PAPRICA should work “out of the box” *except* for the location (and name, if your version differs from mine) of the Silva seed database. The default location in the **paprica_place_it.py** script is inside the genome_finder directory, which may not be where you want to store this. Once you’ve set the locations for everything in the scripts this structure allows you to migrate the python scripts around as needed. **If you do not plan to build a new database from scratch** you can ignore the **paprica_make_ref**, **paprica_build_core_genomes**, and **paprica_parse_blast_xml** scripts.

Next make the **paprica_light.sh** script executable by typing:

```
chmod a+x paprica_light.sh
```

Execute the **paprica_light.sh** script by typing:

```
./paprica_light.sh
```

The script will start by downloading the very large database. This will take a while as the outgoing ftp connection is not as fast as we would like. Make sure that you comment out this and the tar lines for subsequent runs! The script will then attempt to run each of the steps in the PAPRICA pipeline using the provided test.fasta. For your own analysis replace query=test with query=yourfileprefix.

Introduction

PAPRICA is a pipeline to conduct a *metabolic inference* on a collection of 16S rRNA gene sequences. Given a set of 16S rRNA gene sequences the pipeline returns a collection of metabolic pathways that are expected for the observed taxa. The abundance of each metabolic pathway will be corrected for 16S

rRNA gene copy number. PAPRICA as designed with the analysis of large 16S rRNA gene sequence libraries in mind, such as those generated by 454 or Illumina sequencing, but is also appropriate for small datasets.

There are two options for running PAPRICA. If you would like to generate the database from scratch follow the guidelines in the **paprica.sh** script. If you would like to skip the (computationally intensive) database build, you can download a prebuilt version. Use the **paprica_light.sh** script to download and use this database.

Requirements - PAPRICA

PAPRICA requires a Linux-like operating environment. It was developed on Ubuntu. PAPRICA uses several open source tools in addition to some novel components. Most of these are tools that you should probably have if you're working with 16S rRNA gene data anyway, and all come with precompiled executables for Linux/OSX. Required tools that should be located in your path are:

1. [Python 2.7](#)
2. [Biopython](#)
3. The Python [joblib](#) module
4. [Blast+](#) (and the [16SMicrobial](#) database)
5. [Mothur](#) (and one of the Mothur formatted alignment databases, preferably the Silva [seed](#))
6. [FastTreeMP](#)
7. [pplacer \(including Guppy and Taxit\)](#) and [Seqmagik](#) from the Matsen group
8. [Pathway-Tools](#)
9. [Gnu parallel](#)

Requirements – PAPRICA_LIGHT

As above, except only the following are dependencies:

1. [Python 2.7](#)
2. [Mothur](#) (and one of the Mothur formatted alignment databases, preferably the Silva [seed](#))
3. [pplacer \(including Guppy and Taxit\)](#) and [Seqmagik](#) from the Matsen group
4. [Pathway-Tools](#)
5. [Gnu parallel](#)

Components

Seven Python scripts are provided to create the database and execute the metabolic inference. The **paprica.sh** script is a template to guide database construction and analysis. An addition text file, `mean_phi_values.txt`, provides data on the expected genomic plasticity of each genome.

The **paprica_light.sh** script only executes `paprica_place_it.py`, `paprica_get_genomes.py`, and `paprica_tally_pathways`.

Basic operation

The first time **paprica.sh** executes it must construct the database necessary to conduct the metabolic inference. Counterintuitively it needs to run part of an analysis before it can construct the database. Database construction is very time intensive and can take several days, however, subsequent executions are fairly fast.

Near the top of each script is a set of variables that should be modified by the user to reflect directory pathways and other aspects of the operating environment. A brief description of each script follows.

paprica_make_ref.py (only necessary to construct database)

1. Downloads all completed genomes from Genbank using wget
2. Finds the 16S rRNA genes in these genomes using blastn and the 16SMicrobial database
3. Compiles 16S rRNA gene copy number with the provided phi values
4. Outputs:
 - a. Creates, in the reference directory you identify, a directory structure mimicking the relevant portion of the Genbank FTP site.
 - b. combined_16S.fasta – a fasta file of the first 16S rRNA gene identified in each genome.
 - c. genome_data.txt – a tab-delimited file giving each genome, the phi score of genomic plasticity, and the number of 16S rRNA gene copies for that genome.

paprica_parse_blast_xml.py (necessary to construct database)

1. A dependency of **paprica_make_ref.py**.

paprica_place_it.py [query without extension] [reference package name] (necessary to construct database and for analysis)

1. Uses the 16S rRNA genes from the completed genomes to build a reference tree with FastTreeMP after alignment with Mothur
2. Uses taxit to convert the alignment and tree into a reference package
3. Executes pplacer to place query reads on the reference tree
4. Outputs:
 - a. combined_16S.refpkg (directory) - When run to build a reference package (see paprica.sh) produces a directory containing the reference package. This directory will be located in the identified reference directory. Also produces some additional Mothur outputs.
 - b. [query]_combined_16S.pplacer.filter.jplace - When run to execute a phylogenetic placement produces a jplace file that can be passed to Guppy.

paprica_build_core_genomes.py [query.fat – a phyloxml tree] (necessary to construct database)

1. Uses a phyloxml format “fat” tree and blastn to build a consensus genome for each node on the reference tree.

2. Generates statistics for each consensus genome, including size, expected degree of genomic plasticity, and expected 16S rRNA gene copy number.
3. Outputs:
 - a. **combined_16S.core_genomes** (directory) – A directory tree located in the identified reference directory containing the consensus genome for each internal node on the reference tree and some summary statistics.

paprica_get_genomes.py [query.csv – q guppy csv file of placements] [query – a prefix for output]
(necessary for analysis)

1. Uses a csv file of edge placements generated by pplacer and Guppy (you must generate these files manually, see the paprica.sh script) to identify which consensus genomes should be included in the final metabolic inference.
2. Uses gene copy number information to correct the number of times that each consensus genome is represented in the final metabolic inference.
3. Checks to see if a Pathway Genome Database (PGDB) already exists for each included genome.
4. If a PGDB does not exist, it sets up the necessary files for Pathway-Tools and prepares a script to execute the Pathway-Tools pathologic module. The user should execute the generate_pgdb.sh script after this script completes. Depending on the number of PGDBs that need to be created generate_pgdb.sh can take some time to finish. Generally, the more analyses you run, the faster this gets, because fewer new PGDBs need to be created each time.
5. Calculates a confidence score for each sample, based on the expected genomic plasticity of the consensus genome and the relative size of the consensus genome compared to the size of all daughter genomes.
6. Outputs:
 - a. **generate_pgdb.sh** – a shell script in the working directory that must be executed manually to generate the PGDBs.
 - b. Other files in the reference directory necessary for Pathologic execution.
 - c. [query].edge_tally.txt – a tab-delimited file containing the placement edge number, number of placements, number of placements corrected for 16S rRNA gene copy number, size of the core genome, mean clade genome size, standard deviation of the clade genome size, $1 - \phi$, and the 16S rRNA gene copy number.

paprica_tally_pathways.py [query – the input csv file without extension] (necessary for analysis)

1. Uses the information in [query].edge_tally.txt and in the pathway_report.txt file created for each PGDB to generate a list of all the metabolic pathways that are inferred for the sample.
2. Outputs:
 - a. [query].pathway_summary.txt – A tab-delimited file of each pathway and the number of times it is inferred for the sample.
 - b. [query].pathway_detail.txt – A tab-delimited file of each consensus genome, the number of times it occurs (corrected for 16S rRNA gene copy number), and the pathways that were predicted for that genome.

A quick note on directories and workflow

You can run the scripts in whatever directories you choose, provided you set the variables appropriately in each script. I find it easiest, however, to leave the database in the ~/genome_finder directory. If you are building the database from scratch the make_ref, blast_xml, build_core_genomes, and a copy of place_it should be located in this directory. The reference directory is constructed inside this directory and ~/genome_finder/[reference directory]/ serves as the location of the reference directory in the remaining scripts (including the copy of place_it used for analysis). I create a new copy of the remaining scripts (including place_it) in each working directory. The scripts will output to whatever directory they happen to be located in.

Because some initial analysis is required before build_core_genomes can construct the database (it needs a reference tree with node numbers to work from – I can only figure out how to get this from a guppy derived phyloxml tree) a small test.fasta file is provided. In addition to providing the necessary files to build the database this set can be used to test the rest of the pipeline. The sequences are all very similar Actinobacteria reads and should place in the neighborhood of *Clavibacter* (they aren't quite *Clavibacter*, but that's the closest sequenced genome). If they end up somewhere else pplacer is doing something wonky and you will need to investigate the parameters used by mothur and pplacer. The default parameters in the relevant scripts were determined by considerable experimentation and should work fine (see notes below on Technical Details).

Technical Details and Cautions

Phylogenetic placement

A major requirement for pplacer to work correctly is 1) a good reference tree and 2) a high quality alignment. I was warned against using the Silva alignment to build the reference tree on account of its "gappy" natures, however, alignment against greengenes, and de novo alignments with clustalo and MAFFT, produced poor quality trees (and the de novo methods take a very long time). We determined that alignment to the Silva seed produces the best tree and is very fast. A soft mask is used to filter out hypervariable positions, operationally defined as those with < 50 % of sequences having the same base. In my mind that reduces much of the "gappy" nature of the alignment, and *seems* to leave enough positions to construct a good tree.

Ideally the same mask would be applied to the combined query and reference alignment. We found that with longer 454 reads this works okay. With short Illumina reads (which naturally target a hypervariable region) too many positions are masked and there is not enough information left for a good phylogenetic placement.

Our solution is to generate a reference tree using the 50 % masking, and a second tree on the full alignment to provide the mutation model parameters that pplacer needs. This second tree will have "bad" in the file name, only the statistics in the log file are used. Even with this method there are some troublesome placements; the *Bdelle vibrio* clade in particular receives incorrect placements unless

excessive masking is applied. Lacking a better solution we have eliminated those genomes (and the corresponding 16S rRNA genes) from the database.

Extensive testing has not identified any other problem areas of the tree, **but I suggest spot checking your dataset to make sure that the phylogenetic placement portion of the analysis is working correctly.** The easiest way to do this is to grab a handful of sequence ID's corresponding to an interesting or suspect placement from the .csv file produced by Guppy, extract those sequences from your original fasta file, and blast them using the online NCBI server. Blast is a rough analysis, but if it returns taxa that approximate the region of placement on the reference tree than you should be good to go!

Related Methods

PICRUSt - <http://picrust.github.io/picrust/>

Development Objectives

- Standardize inputs and outputs – some scripts require extension, some do not. All input files should require extension, output name arguments should not.

IN PROGRESS